

Data 2020: Final Project

Introduction

Given the dataset, there are many ways to build multilevel models. I experimented with a few different models, mostly the results were discouraging — but I am confident there is a lot of room for improvement. In the following paragraphs I will describe my process and my findings.

Selecting a level

After the discussions last class I decided it made sense to use County and State as levels. There were 47 states in our dataset with recorded police killings and 238 counties. Intuitively, it seemed that the states would yield more informative results simply because there were more recorded killings in each state. With that said, I could imagine there being more variance in the factors related to police killings across a state as opposed to a county. This could be a potential downside.

It would also be interesting to use the race of the deceased as a level and train the model to predict something like census tract per capita income, or unemployment. Or, we could go the other way, and split the tracts into quintiles based on per capita income or unemployment and, using the quintiles as levels, try to predict the race of the deceased. In my exploratory analysis report I mentioned that I was worried about arriving at any conclusions in my inference of the data. Specifically, I thought that income could be a confounding factor. This could be a way to control for that (of course there are possible routes to explore this that would not involve multilevel models).

Another interesting level to consider would be the “armed” variable. It is easy to imagine there would be different slopes or intercepts based on this value of this category.

As of now, I have made models with state and county as levels.

Pre-Processing The Data

This proved more time consuming than anticipated. No surprise there. More than anything, I had to figure out a way to aggregate the census data for each state and county. Below is the code for how I did this for each state. The process for aggregating the county level data was the same.

```
28
29 toytableState = acs %>%
30   group_by(State) %>%
31   summarise(TotalState = sum(TotalPop),
32             StateMen = sum(Men)/TotalState,
33             #TotalWomenState = sum(Women),
34             StateUnemployment = (sum(TotalPop*(Unemployment*.01)))/TotalState,
35             StateIncomePerCap = sum(TotalPop*IncomePerCap)/TotalState,
36             StatePoverty = (sum(TotalPop*(Poverty*.01)))/TotalState,
37             StateFracHispanic = (sum(TotalPop*(Hispanic*.01)))/TotalState,
38             StateFracBlack = (sum(TotalPop*(Black*.01)))/TotalState,
39             StateFracWhite = (sum(TotalPop*(White*.01)))/TotalState)
40
```

Building a Model

This was fairly trivial. One thing to note was that there weren't enough observations to build a county level varying slope model, which is unfortunate, because it would have been interesting to compare the county level varying slope model to the state level varying slope model. It's also validating, remember I voiced my concern about the number of counties with recorded police killings in the selecting a level section. In addition, the model failed to converge when I tried to give the armed feature a varying slope in the state level model.

```
85 InterceptStateModel <- glmer(Minority ~ age + gender + armed + (1|State),
86                             family = binomial("logit"), data = datastateuse)
87
88 #unfortunately the model fails to converge for features other than age and
89 #gender having varying coefficients
90 OneSlopeStateModel <- glmer(Minority ~ age + gender + armed + (1+age|State),
91                             family = binomial(link = "logit"), data = datastateuse)
92 TwoSlopeStateModel <- glmer(Minority ~ age + gender + armed + (1+age+gender|State),
93                             family = binomial(link = "logit"), data = datastateuse)
94
95 InterceptCountyModel <- glmer(Minority ~ age + gender + armed + (1|County),
96                              family = binomial(link = "logit"), data = datacountyuse)
97
98 #Unfortunately, we don't have enough observations to run a varying slope model for the county
99 #SlopeCountyModel <- glmer(Minority ~ age + gender + (1+age+gender|County),
100 #                           family = binomial("logit"), data = datacountyuse)
101
```

Model Assessment

The results were nothing to write home about. When I compared the three state models together the AIC and BIC scores varied marginally corresponding to the model degrees of freedom. The likelihood values were all the same as were the deviances. Below you can see the

results just described, obtained using the “anova” function. The chisq value for the varying slope model where the age feature had a varying coefficient was significant.

```
> ## Compare models
> kable(anova(InterceptStateModel, OneSlopeStateModel, TwoSlopeStateModel), type = "pandoc", caption = "Table 1: Comparison of varying/intercept slope and varying intercept on State")
```

	Df	AIC	BIC	logLik	deviance	Chisq	Chi	Df	Pr(>Chisq)
InterceptStateModel	5	551.3337	571.5469	-270.6668	541.3337	NA	NA		NA
OneSlopeStateModel	7	555.3030	583.6015	-270.6515	541.3030	0.0306642	2		0.9847848
TwoSlopeStateModel	10	560.9640	601.3903	-270.4820	540.9640	0.3390739	3		0.9525197

Below are some diagnostics obtained from the county level model. They are similar to the ones above though the AIC and BIC numbers are smaller (there are 5 degrees of freedom, hence `df.resid = 416`).

AIC	BIC	logLik	deviance	df.resid
547.5	567.7	-268.7	537.5	416

Conclusions

As I mentioned there is still much to do! Obviously, my results leave something to be desired. Additionally, I would like to run more diagnostics to get a better sense of just how bad the models are. Nonetheless, I successfully built a few multilevel models, something that had eluded me up until this point and I feel like I have the start of a foundation to build on. I look forward to building different multilevel models and improving the ones I have.¹

¹ I have not included my code for this section since it was all fairly straightforward and boilerplate.